

Using risk and context maturity to decide the depth of AI use

A short briefing on choosing between human-led, agent-led, or autonomous work

Dorian Vocolka · September 2026

1. Categorizing the cooperation spectrum

This note is a short briefing on how much of a piece of work a human should still own, and how much an agent may run. The common description is a spectrum: from work done entirely by people to work done entirely by agents.

Agents here are LLM-enabled programs given a goal and some tools. They can take a sequence of steps toward that goal without a person directing each step.

McKinsey, for example, published an article that names three of those stops for AEC workflows [\[1\]](#). When to use each mode is the job of the briefing that follows, not of the labels themselves.

- **Human-led, agent-enabled.** A person owns the work end to end, usually talking to other people, and uses agents for search, drafts, checks, and gathering information. The person still decides.
- **Agent-led, human in the loop.** An agent runs the main loop: gather facts, compare options, estimate impact, draft a recommendation. A person stays in the loop for the calls that need judgment and that someone will be held to.
- **Autonomous.** The agent runs the workflow end to end. People are used mainly to design the workflow and to oversee it. A human step is an exception, not a required station.

2. Decision lens: risk + context

A fast way to choose a mode is to score the task on two axes:

- **Risk:** What is the cost of a wrong or incomplete result? (rework, safety, legal, reputation, system-wide debt)
- **Context maturity:** Is the needed knowledge already explicit, structured, and loadable—or is it scattered in people's heads, brownfield systems, and tribal process?

Working rule

- High risk and/or low context maturity → human-led, agent-enabled.
- Medium risk with mostly available context → agent-led, human in the loop.
- Low risk with clear specs and cheap validation → autonomous.




Figure 1. Decision lens: risk and context maturity.

Steps to follow when expanding or improving the context is not an option (fast action required):

1. Classify the next task as high / medium / low risk before choosing tools.
2. If context maturity is low (people, systems, history not loadable), do not push into autonomous mode.
3. Put human checkpoints where mistakes are expensive; put autonomy where mistakes are local and testable.
4. Expect the threshold to move: more tasks will become medium or low risk as models and captured context improve. Re-score periodically; do not freeze today's boundary as permanent.

3. Use the objective method to avoid widening the gap

In 2025 and 2026 it was memed among people who write software that “the hard part was never the coding,” as generation got cheaper. The hard part is a specification good enough to start. The mirror is also true: for people who speak for the customer, the hard part was never listing market needs.

Between those two groups sits work that only exists if both sides do it: turning a market need into a specification that can be built, turning a technical constraint into something a customer can accept and doing it the right way (e.g. scalable). That joining work is expensive—meetings, arguments, failed passes, people. Teams may try to skip it with a one-time handover: a form filled out “just so,” reviewed once, then built. When each side stays in its own half, the joining work is simply not done. Call that missing stretch the gap.

Engineers must take on the other half: what is needed, which requirements matter, what purpose the system is for.

Sales and customer representatives must take on the other half: constraints, coupling, what a change really costs.

Neither side is trained, equipped, or usually allowed to do that well, so the gap stays. An agent can look as if it removes the need for that joining work: a non-technical person asks a technical question and gets a fluent reply, or an engineer asks what the customer needs and gets a fluent reply. The reply does not do the joining work. That is the reason to judge risk and context on the work itself first, and only then decide how much of it an agent should run.

4. Quick mapping table and forward-deployed engineering example

Risk band	Typical work	Preferred mode	Why
High	System design, brownfield “what touches what”, legal/safety requirements, stakeholder “what they actually need”	Human-led, agent-enabled	Value sits in judgment and tacit context; wrong call is expensive
Medium	Local technical choices once architecture exists; structured research; draft proposals	Agent-led, human in the loop	Agent does the volume; human owns exceptions and final call
Low	Implementation inside a decided component; mechanical transforms; routine cleanup with clear rules	Autonomous	Failure is cheap to detect and redo; debt stays local

Same lens on one kind of delivery job—forward-deployed engineering style work:

- Requirements & system fit: Human-led, agent-enabled. Agents help interview prep, note synthesis, and option comparison—not final ownership of “what we should build.”
- Stakeholder alignment: Human-led or agent-led, human in the loop for narrative and commitments.
- Coding inside agreed boundaries: Often low-to-medium risk → agent-led, human in the loop, or autonomous with tests and review gates.
- Brownfield “which components?”: High risk today; context is not orderly. Human-led, agent-enabled until system-specific knowledge is captured.

5. In short

An agent that answers fluently does not replace the joining work between people who know the purpose and people who know the system. Score the piece of work on risk and on whether the needed context is already loadable. Choose human-led and agent-enabled, agent-led with a human in the loop, or autonomous from that score—not from the tool.

Use this when a decision has to be made quickly. The method itself is a small effort—20 percent in the Pareto sense—and captures most of the benefit, if the current habit is “automate everything” or a gut call. That is still the usual starting point in late 2026.

6. References

- [1] Daniel Ahmoye, Erik Sjödin, and Jose Luis Blanco, with Aleem Mawji, Marita Winslade, and Patrick Rogers (McKinsey & Company), “How AI is reshaping the future of the AEC industry” (15 July 2026). <https://www.mckinsey.com/industries/engineering-construction-and-building-materials/our-insights/how-ai-is-reshaping-the-future-of-the-aec-industry>

7. Publication notes

Cite. Vocolka, D. (2026). Using risk and context maturity to decide the depth of AI use: A short briefing on choosing between human-led, agent-led, or autonomous work. September 2026.

Independence. Written in a personal capacity. It does not represent the views of any employer. No proprietary or non-public work product is used. Sources are public.

AI assistance. Drafting and editing were assisted by large language model tools (including Grok). Framing, source selection, judgments, and the final text are the author's responsibility.

License. Creative Commons Attribution–NonCommercial 4.0 International (CC BY-NC 4.0). Share with attribution for non-commercial use.

Contact. @DorianVocolka